

CEVE 543 Final Project

Linear and Non-Linear Dimensionality Reduction in Climate Science: Methods and Applications of PCA and Autoencoders

Erika Adjei

2025-11-14

1 Introduction: What is Dimensionality Reduction?

The size and scale of high dimensional data is vast and noisy. This is particularly true in the field of climate science. Within the field of climate science, dimensionality reduction methods can be used to simplify massively complex datasets, identify major patterns of variability, and improve the accuracy and interpretability of climate models. Specifically, this helps to constrain/differentiate between climate noise, internal variability signals, and human-driven climate responses, allowing climate scientists to, for example, quantify modes of climate variability like the El Niño–Southern Oscillation (ENSO)(DelSole et al. 2022).

2 Principal Component Analysis (PCA)

3 What is PCA?

PCA, or Principal Component Analysis, is a dimensionality reduction technique that reduces the dimensionality of data by transforming it into a new set of orthogonal axes (called principal components) that capture the most variance. This technique projects data onto these components in a way that maximizes the variance in fewer dimensions (DelSole et al. 2022).

4 How does PCA work

First, standardize your data by normalizing (subtracting the mean and dividing by standard deviation for each feature to ensure that feature has equal weight for PCs). Then, compute the covariance matrix to identify how variables relate to each other. Calculate eigenvectors

and eigenvalues of your covariance matrix to determine the direction and strength of variance in your data. Finally, project your data onto top principal components (DelSole et al. 2022).

5 Strengths, Limitations, and Appropriate Use Cases of PCA:

For small to moderately-sized, linearly separable datasets, PCA is an effective, computationally efficient method that provides an intuitive understanding of variance and how features relate to each other. Further, the eigenvalue decomposition and modern linear algebra libraries that serve as the mathematical foundation for this technique are well-studied. Disadvantages of this technique include the usage of variance, PCA's sensitivity to scale, and the linear assumptions this technique makes. PCA assumes that the directions of maximum variance are most important, and this does not hold true in certain situations. Additionally, because PCA is sensitive to the scale of input variables, variables with larger numerical ranges will dominate PCs. Standardization (usually mean centering and scaling to unit variance) can address this problem, but this can change the interpretation of results. PCA also only captures linear relationships between variables. So, if your data has non-linear patterns, PCA may not capture the complexity of variable interactions. Thus, PCAs should be used when data variable dimensionality is high and multi-collinearity exists between variables (Idrees 2021)(DelSole et al. 2022).

6 PCA Mathematical Framework (DelSole et al. 2022)

7 Key Assumptions

This approach assumes that your data is discretized, a linear combination of orthogonal basis vectors, and mean centered.

Consider a data set $(Y(t,s))$ with

- $(t = 1,...,N)$ time indices,
- $(s = 1,...,M)$ spatial indices.

We can write:

$$Y(t, s) = f(t)e(s).$$

In matrix form, the data is an $(N \times M)$ matrix (Y) ,
with each row a time step and each column a spatial location.

8 Goal of PCA

Find a low-rank approximation of (Y) using (K) modes:

$$Y \approx Y_K = f_1 e_1^T + f_2 e_2^T + \dots + f_K e_K^T$$

9 PCA Optimization Problem

The first principal component solves:

$$\min_{f_1, e_1} \|Y - f_1 e_1^T\|_F^2.$$

After removing the first component, the PCA is applied to the residual:

$$Y - f_1 e_1^T.$$

10 Solution via Singular Value Decomposition (SVD)

PCA is equivalent to performing the SVD:

$$Y = USV^T,$$

where

- (U) contains left singular vectors,
- (V) contains right singular vectors,
- (S) is a diagonal matrix with entries (s_1, ..., s_R) (the singular values).

The rank-(K) approximation is:

$$Y_K = s_1 u_1 v_1^T + s_2 u_2 v_2^T + \dots + s_K u_K v_K^T.$$

11 Identifying PCs and EOFs

Define:

$$f_k = s_k u_k, \quad e_k = v_k.$$

Thus:

$$Y_K = f_1 e_1^T + \cdots + f_K e_K^T.$$

Here,

- (e_k) are Empirical Orthogonal Functions (EOFs),
 - (f_k) are Principal Components (PCs).
-

12 Variance Explained

Because singular vectors are orthogonal:

$$\|Y\|_F^2 = s_1^2 + s_2^2 + \cdots + s_R^2.$$

Fraction of variance explained by the first (K) modes:

$$\frac{s_1^2 + s_2^2 + \cdots + s_K^2}{s_1^2 + s_2^2 + \cdots + s_R^2}.$$

13 Example of PCA:




Figure 12.1 (a) Leading EOF of the December SST from the Extended Reconstructed SST of Huang et al. (2015). The percent variance explained by the leading component is indicated in the title. The value at each grid point gives the departure from climatology in units of degrees Kelvin per unit PC time series. (b) The corresponding leading PC time series.

Figure 1: Example of a PCA output (Huang et al. 2015)(DelSole et al. 2022). Shows leading EOF explaining global Decemebr SSTs .

14 What are Autoencoders:

Autoencoders are a type of neural network that compresses high-dimensional, non-linear data into a lower-dimensional space, then reconstructs, or decodes, the original data from this compressed representation. Input and output dimensions are kept the same. This ML technique captures hidden or random variables that, despite not being directly observable, fundamentally inform the way data is distributed (Gómez et al. 2004).

15 How do Autoencoders work?

The first part of an Autoencoder framework compresses input data into a lower dimensional latent space called a bottleneck. This aims to capture the initial data's most important features in fewer dimensions. The compressed version of this data is represented in the bottleneck layer, and this has fewer neurons than the input. The decoder, or the second part of the Autoencoder network, attempts to reconstruct the original input data from the compressed latent space (Myrzaliyeva 2022).

16 Strengths, Limitations, and Appropriate Use Cases of Autoencoders:

Some advantages of Autoencoders include their ability to reduce the dimensionality of complex, non-linear datasets (ex: images and text) where linear methods like PCA may fail, and the flexibility and vast applicability of neural network architecture are used for feature extraction, like data compression, image denoising, anomaly detection and facial recognition. Disadvantage-wise, Autoencoders are computationally more expensive and require more data to train effectively. Additionally, the features learned in the latent space (results of Autoencoders) can be hard to interpret. Further, as with many deep learning techniques, Autoencoders suffer from the black box problem. Autoencoders should be used when dealing with very large, non-linear datasets with abstract bounds (ex: images, pictures, etc.) (Idrees 2021) (Myrzaliyeva 2022).

17 Autoencoder Mathematical Framework (Gómez et al. 2004) (Ross 2009):

18 Key Assumptions

Input features must be scaled to a consistent range using techniques like Z-Score normalization or Min-Max Scaling (Myrzaliyeva 2022).

19 Autoencoder Mathematical Framework

19.1 Encoder

The encoder maps an input vector $x \in \mathbb{R}^d$ to a latent representation $z \in \mathbb{R}^k$:

$$z = f(x; \theta_e).$$

For a single-layer encoder:

$$z = \sigma_e(W_e x + b_e).$$

19.2 Decoder

The decoder reconstructs the input from the latent representation:

$$\hat{x} = g(z; \theta_d).$$

For a single-layer decoder:

$$\hat{x} = \sigma_d(W_d z + b_d).$$

19.3 Full Autoencoder Mapping

The full mapping is:

$$\hat{x} = (g \circ f)(x) = g(f(x; \theta_e); \theta_d).$$

19.4 Reconstruction Loss

The autoencoder is trained by minimizing reconstruction error:

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N L(x^{(i)}, \hat{x}^{(i)}) = \frac{1}{N} \sum_{i=1}^N L(x^{(i)}, g(f(x^{(i)}; \theta_e); \theta_d)).$$

19.5 Common Loss Functions

Mean Squared Error (MSE):

$$L_{\text{MSE}}(x, \hat{x}) = \frac{1}{d} \sum_{j=1}^d (x_j - \hat{x}_j)^2.$$

Binary Cross-Entropy (BCE):

$$L_{\text{BCE}}(x, \hat{x}) = -\frac{1}{d} \sum_{j=1}^d [x_j \log(\hat{x}_j) + (1 - x_j) \log(1 - \hat{x}_j)].$$

19.6 Optimization

Autoencoder training solves:

$$\theta^* = \arg \min_{\theta} J(\theta).$$

Typical gradient-descent update:

$$\theta \leftarrow \theta - \eta \nabla_{\theta} J(\theta).$$

20 Example of Autoencoder




Figure 3.7: Original space and latent space based clusters.

Figure 2: Example of an Autoencoder output (Myrzaliyeva 2022). Latent space output is being compared to initial input data.

21 Application of Dimensionality Reduction to Class:

CEVE 543's course work centers around quantifying hydroclimate extremes, their uncertainties, and the related impacts on relevant communities using varying statistical methods. Within the context of climate change, Autoencoders and PCAs accomplish this goal. By collapsing noisy, high-dimensional data into low-dimensional modes of variance, these tools allow for the attribution and forecasting of hydroclimate extremes. This informs underlying patterns (ex: modes of internal variability like AMO and ENSO) driving pluvials, droughts, seasonal rainfall patterns, and regional rainfall extremes. Additionally, the use of dimensionality reduction tools like PCAs and Autoencoders underpin the statistical framework of modern-day climate reanalyses (ex: LMR, PHYDA), and improves the use of statistical tools like stochastic weather generators.

22 Work Cited:

DelSole T, Tippet M. Principal Component Analysis. In: Statistical Methods for Climate Scientists. Cambridge University Press; 2022:273-297.

Gómez, A. J., Zhou, C. S., Timmermann, A., and Kurths, J.: Nonlinear dimensionality reduction in climate data, *Nonlin. Processes Geophys.*, 11, 393–398, <https://doi.org/10.5194/npg-11-393-2004>, 2004.

Idrees, H. Autoencoders vs PCA: Dimensionality Reduction for Complex Data. Medium; [https://medium.com @ hassaanidrees7 autoencoders-vs-pca-dimensionality-reduction-for-complex-data-e07d4612b711](https://medium.com/@hassaanidrees7/autoencoders-vs-pca-dimensionality-reduction-for-complex-data-e07d4612b711) (2024)

Myrzaliyeva, M. : Investigating the Impact of Climate Change on Weather Regimes Using Dimensionality Reduction with Deep Autoencoders, Vrije Universiteit Brussel. https://cris.vub.be/ws/portalfiles/portal/94889169/MA_ACS_Myrzaliyeva_Madina_S3_2122_final.pdf, 2022:1-41. . Ross, Ian. “Nonlinear Dimensionality Reduction Methods in Climate Data Analysis.” *arXiv: Atmospheric and Oceanic Physics* (2009): n. pag.